

**Title:** Apache Kafka Interview Questions and Answers (100+ Q&A)

**Introduction:** This guide contains over 100 interview-level questions and answers on Apache Kafka, covering Zookeeper, KRaft mode, producers, consumers, transactions, and integration with ELK and Splunk. It is intended as a comprehensive preparation resource for software engineers, DevOps, and data engineers.

---

## 1. Kafka Basics

### 1. What is Apache Kafka?

2. Answer: Kafka is a distributed streaming platform used for building real-time data pipelines and streaming applications.

### 3. What are the main components of Kafka?

4. Answer: Producer, Consumer, Broker, Topic, Partition, Zookeeper (optional in KRaft).

### 5. What is a Kafka topic?

6. Answer: A topic is a named logical channel to which messages are published by producers.

### 7. What is a partition?

8. Answer: A partition is a subset of a topic; partitions allow for parallelism and scalability.

### 9. Difference between Kafka and traditional message queues?

10. Answer: Kafka provides persistent storage, high throughput, horizontal scalability, and message replay, unlike traditional queues.

### 11. Role of producer?

12. Answer: Sends messages to Kafka topics.

### 13. Role of consumer?

14. Answer: Reads messages from topics and processes them.

### 15. Push vs pull consumer model?

16. Answer: Kafka uses a pull model where consumers request batches of messages.

### 17. Kafka retention?

18. Answer: Kafka retains messages for a configured time or size regardless of consumption.

19. **Durable vs ephemeral messaging?**

20. Answer: Kafka messages are durable; traditional queues may discard messages after delivery.

## 2. Kafka Architecture

1. **Explain Kafka broker.**

2. Answer: Broker is a server that stores messages and serves client requests.

3. **Kafka cluster?**

4. Answer: A set of brokers working together to provide scalability and fault tolerance.

5. **How does Kafka achieve scalability?**

6. Answer: By partitioning topics and distributing partitions across brokers.

7. **Leader and follower replicas?**

8. Answer: Leader handles reads/writes; followers replicate leader data.

9. **ISR (In-Sync Replica)?**

10. Answer: Set of replicas that are fully caught up with the leader.

11. **Failover handling?**

12. Answer: If a leader fails, a follower from ISR becomes leader.

13. **Role of Zookeeper?**

14. Answer: Maintains cluster metadata, broker status, leader election.

15. **Topics and partitions purpose?**

16. Answer: For organizing and scaling data streams.

17. **Replication factor vs partitions?**

18. Answer: Replication factor is number of copies; partitions enable parallelism.

19. **High throughput?**

20. Answer: Achieved via batching, compression, and sequential disk writes.

## 3. Kafka Producers and Consumers

1. **Producer acknowledgments (acks)?**

2. Answer: 0=fire-and-forget, 1=leader ack, all=all ISR ack.

**3. Sync vs async producer?**

4. Answer: Sync waits for ack; async sends without blocking.

**5. Idempotent producer?**

6. Answer: Ensures messages are not duplicated on retries.

**7. Consumer groups?**

8. Answer: Multiple consumers share topic partitions for parallel processing.

**9. Consumer offsets?**

10. Answer: Track last read position in partitions.

**11. Auto vs manual commit?**

12. Answer: Auto commits offsets automatically; manual allows controlled commit.

**13. Rebalancing?**

14. Answer: Redistributes partitions among consumers when group membership changes.

**15. Message key usage?**

16. Answer: Determines which partition a message is sent to.

**17. Message ordering?**

18. Answer: Preserved per partition.

**19. Backpressure handling?**

20. Answer: Consumers poll at their pace; producers can buffer messages.

## **4. Kafka Brokers, Partitions, and Replication**

**1. Data replication?**

2. Answer: Leader replicates messages to followers.

**3. Leader election for partitions?**

4. Answer: Zookeeper (or KRaft) selects leader among replicas.

**5. Broker metadata management?**

6. Answer: Metadata includes broker IDs, topics, partitions, leader info.

**7. Log segment?**

8. Answer: Segment is a chunk of partition logs stored on disk.

**9. Log compaction?**

10. Answer: Retains only latest value per key.

**11. Parallelism via partitions?**

12. Answer: Consumers can read different partitions concurrently.

**13. Default replication factor?**

14. Answer: Typically 1 (configurable).

**15. Broker failure handling?**

16. Answer: ISR followers take over leader responsibilities.

**17. Data durability?**

18. Answer: Achieved via replication and disk storage.

**19. Committed vs uncommitted messages?**

20. Answer: Committed messages are acknowledged by required replicas.

## **5. Kafka Zookeeper Role**

**1. What is Zookeeper?**

2. Answer: Centralized service for configuration, coordination, and synchronization.

**3. Why Kafka used Zookeeper?**

4. Answer: To manage metadata, leader election, and cluster state.

**5. Metadata stored?**

6. Answer: Broker info, topic configs, partition leaders, ISR info.

**7. Leader elections?**

8. Answer: Zookeeper selects partition leaders when needed.

**9. Broker discovery?**

10. Answer: Brokers register with Zookeeper; clients query Zookeeper.

**11. Broker health monitoring?**

12. Answer: Zookeeper watches broker heartbeat nodes.

**13. Zookeeper vs KRaft?**

14. Answer: KRaft removes Zookeeper dependency, uses Raft internally.

**15. Topic metadata in Zookeeper mode?**

16. Answer: Stored in Zookeeper znodes.

**17. Limitations?**

18. Answer: External dependency, complex deployment, potential performance bottleneck.

**19. Securing Zookeeper?**

20. Answer: Use ACLs, TLS, authentication mechanisms.

## 6. Kafka KRaft Mode

**1. What is KRaft?**

2. Answer: Kafka Raft mode; manages metadata without Zookeeper.

**3. Why no Zookeeper?**

4. Answer: Raft quorum handles leader election and metadata internally.

**5. Kafka controller?**

6. Answer: Node responsible for cluster metadata and leadership.

**7. Metadata storage?**

8. Answer: Stored in internal topics ( `__cluster_metadata` ).

**9. Raft consensus?**

10. Answer: Ensures metadata replication across controllers.

**11. Advantages over Zookeeper?**

12. Answer: Simpler architecture, faster leader election, no external dependency.

**13. Migration from Zookeeper?**

14. Answer: Supported via migration tools in newer Kafka versions.

**15. Leader election in KRaft?**

16. Answer: Raft controller quorum elects leaders.

**17. Quorum controllers?**

18. Answer: Multiple controllers to maintain cluster consensus.

**19. Kafka versions supporting KRaft?**

20. Answer: Kafka 2.8+ with incremental improvements; production-ready 3.3+.

## **7. Kafka Transactions & Exactly-Once Semantics**

**1. Exactly-once semantics (EOS)?**

2. Answer: Guarantees a message is processed exactly once from producer to consumer.

**3. How EOS achieved?**

4. Answer: Idempotent producers + transactional APIs.

**5. Idempotent producer?**

6. Answer: Prevents duplicate messages on retries.

**7. Transaction initiation?**

8. Answer: Producer begins transaction via API.

**9. Transaction IDs?**

10. Answer: Unique ID to identify and manage transactions.

**11. EOS vs at-least-once?**

12. Answer: EOS ensures single delivery; at-least-once may duplicate.

**13. Consumers handling transactional messages?**

14. Answer: Read only committed transactional messages.

**15. Broker failure during transaction?**

16. Answer: Transaction is aborted if not committed.

**17. Producer retries?**

18. Answer: Handled idempotently with transaction coordinator.

19. **Transactional timeouts?**

20. Answer: Configured via `transaction.timeout.ms`.

## 8. Kafka Integration with ELK and Splunk

1. **Kafka → ELK integration?**

2. Answer: Logstash consumes Kafka topics, transforms, sends to Elasticsearch.

3. **Logstash role?**

4. Answer: Parsing, filtering, and enriching Kafka messages before indexing.

5. **Send Kafka messages to Elasticsearch?**

6. Answer: Use Kafka input plugin in Logstash and Elasticsearch output plugin.

7. **Visualize in Kibana?**

8. Answer: Create dashboards on indexed topics.

9. **Kafka → Splunk integration?**

10. Answer: Use Splunk Connect for Kafka or push via HEC.

11. **Splunk HEC?**

12. Answer: HTTP Event Collector for ingesting JSON logs.

13. **Kafka consumers to HEC?**

14. Answer: Write consumer app or use connector to send messages to Splunk HEC.

15. **Kafka → ELK vs Kafka → Splunk?**

16. Answer: ELK supports filtering, dashboards; Splunk supports alerting, SPL queries.

17. **High throughput handling?**

18. Answer: Batch messages, compress,